

# Shortest path

This problem consists of three subtasks: the Bellman-Ford algorithm, the Dijkstra's algorithm, and the difference constraints problem.

Please go to [Piazza Resources](#) to download the related materials first.

## The `Graph` class

`graph.hpp` is a header file defining a `Graph` class that represents a **directed, edge-weighted** graph  $G = (V, E)$  using adjacency lists. For simplicity, we assume that  $V = \{0, 1, \dots, n - 1\}$  and that the edge weights are integers. Considering that you might have not written any graph-related code before, we have implemented two basic functions for you as demonstration:

- The Breadth-First Search algorithm, which has the following signature

```
void bfs(VertexID start, std::invocable<VertexID> auto callback) const;
```

`start` is the id of the vertex where the traversal begins. `callback` is a callable object that is called every time a vertex is visited. For example, the following code (see `examples/graph_example.cpp`) performs BFS on the graph `graph` and prints the id of the visited vertices in order.

```
graph.bfs(0, [](auto x) { std::cout << x << ' '; });
```

- The naive  $O(V^2)$  Dijkstra's algorithm, which has the following signature

```
std::vector<Weight> dijkstra(VertexID source) const;
```

`source` is the id of the source vertex. This function returns a `std::vector<Weight>`, in which the element indexed `i` is the length of the shortest path from `source` to the vertex `i`, or `Graph::infinity` if no path exists.

The behavior is undefined if there exists an edge with negative weight.

By saying "the behavior is undefined", we mean that **you don't need to care about this case**. Do not waste your effort detecting it.

The implementation of this function is in `graph.cpp`, which you need to overwrite (See [below](#)).

Read through the code we have provided in `graph.cpp`, `graph.hpp` and `examples/graph_example.cpp` and make sure you have understood the design. To compile and run `graph_example.cpp`, refer to [compile and run](#).

**Note:** Your submission will not contain the contents in `graph.hpp`. Any modification to that file (e.g. adding a member to the class `Graph`) will have no effect.

## Subtask 1: The Bellman-Ford algorithm

First, you need to implement the Bellman-Ford algorithm that calculates the length of the shortest path from a given source vertex to each vertex, or reports that there is a negative cycle reachable from `source`. It is a member function of `Graph` that has the following signature

```
std::optional<std::vector<Weight>> bellmanFord(VertexID source) const;
```

`std::optional<T>` is a special standard library type that represents **either an object of type `T`, or nothing (represented by `std::nullopt`)**. You can refer to `examples/equation.cpp` to see an example of using it.

This function should return `std::nullopt` if the graph contains a negative cycle reachable from `source`. Otherwise, it should return a vector of numbers, in which the element indexed `i` is the length of the shortest path from `source` to the vertex `i`, or `Graph::infinity` if no path exists. This function should run in  $O(VE)$  time with  $O(V)$  extra space.

Write your implementation in `graph.cpp`.

## Subtask 2: The Dijkstra's algorithm

In `graph.cpp`, there is an implementation of the naive Dijkstra's algorithm whose time complexity is  $O(V^2)$ . Suppose now the graph is not so dense, or more specifically,  $E = o(V^2 / \log V)$ . Overwrite this function to make it run within  $O(E \log V)$  time, with  $O(V)$  extra space.

You may need `std::priority_queue`. We have provided two examples of using it: `examples/priqueue.cpp` and `examples/huffman.cpp`. The latter is a program that computes the Huffman codes given the frequencies of characters.

By default, `std::priority_queue<T>` uses `std::less<T>` (which is `operator<`) to compare elements and models a **max-heap**. If you need a min-heap, just pass `std::greater<>` as the third template argument (see details in the examples).

We have seen some inexperienced students overload `operator<` to do the work that should have been done by `operator>` (e.g. `return lhs.something > rhs.something;`), or pass negated values to the max-heap to *fake* a min-heap. **You should never do that.**

## Subtask 3: System of difference constraints

In a **system of difference constraints** (or, **difference constraints problem**), there are  $n$  variables  $x_0, x_1, \dots, x_{n-1}$  and  $m$  constraints on these variables, where the  $k$ -th constraint ( $0 \leq k < m$ ) is of the form

$$x_{u_k} - x_{v_k} \leq c_k$$

where  $0 \leq u_k, v_k < n$ , and  $u_k \neq v_k$ . We want to find a solution to the given difference constraints problem, that is, to assign the variables with values so that all the constraints are satisfied.

For example, the following is a system of difference constraints with 5 variables and 8 constraints:

$$\begin{aligned}
x_0 - x_1 &\leq 0, \\
x_0 - x_4 &\leq -1, \\
x_1 - x_4 &\leq 1, \\
x_2 - x_0 &\leq 5, \\
x_3 - x_0 &\leq 4, \\
x_3 - x_2 &\leq -1, \\
x_4 - x_2 &\leq -3, \\
x_4 - x_3 &\leq -3.
\end{aligned}$$

One solution to this problem is  $(x_0, x_1, x_2, x_3, x_4) = (-5, -3, 0, -1, -4)$ , which you can verify directly by checking each inequality. Note that a system of difference constraints usually have infinitely many solutions, but some systems may have no solutions at all. The following is a system with 3 variables and 3 constraints that have no solutions:

$$\begin{aligned}
x_0 - x_1 &\leq -1, \\
x_1 - x_2 &\leq -1, \\
x_2 - x_0 &\leq -1.
\end{aligned}$$

A system of difference constraints can be modeled as a graph, and then solved by solving a single-source shortest-path problem on the graph. **Read Section 24.4 of *Introduction to Algorithms, Third Edition*** which is available in Piazza Resources.

The difference constraints problem that you need to solve will be presented as the class `Problem` in `problem.hpp`. Read through the code in `problem.hpp` and understand the structure of that class. Note that we have implemented some useful functions for you, such as `hasNegativeConstant()`, `getNumVars()` and `getConstraints()`.

Your task is to implement the `solve` function in `main.cpp`, which has the signature

```
std::optional<Problem::Solution> solve(const Problem &problem);
```

It solves the given difference constraints problem `problem`, and returns the solution to it or `std::nullopt` if it has none. A solution is of type `Problem::Solution` i.e. `std::vector<Problem::Value>` in which the element indexed `i` is the value assigned to the variable  $x_i$ .

For testcase 1~5,  $c_k$  may be negative, and your algorithm should run within  $O(n^2 + nm)$  time. For testcase 6~10,  $c_k$  are non-negative, and your algorithm should run within  $O((n + m) \log n)$  time.

## Compile and run

### Basic knowledge

Read [CompileBasics.md](#) if you need.

### Compile multiple files

The functions `Graph::dijkstra` and `Graph::bellmanFord` are defined in `graph.cpp`. Therefore, every time you compile a program that calls those functions, `graph.cpp` should be compiled and linked as well. You will see an error like

```
/usr/bin/ld: /tmp/ccxICQ7U.o: in function `main':
graph_example.cpp:(.text+0x15e): undefined reference to `Graph::dijkstra(unsigned long) const'
collect2: error: ld returned 1 exit status
```

if you did not link them correctly. The simplest way is to also pass the path to `graph.cpp` as an argument to the compiler. For example, to compile `solve.cpp`, `cd` to the directory `attachments` first, and then

```
g++ solve.cpp graph.cpp -o solve -std=c++20
```

To compile `examples/graph_example.cpp`: If the current working directory is `attachments`,

```
g++ examples/graph_example.cpp graph.cpp -o examples/graph_example -std=c++20
```

This will generate an executable named `graph_example` (`graph_example.exe` on Windows) in the directory `attachments/examples`. If the current working directory is `attachments/examples`,

```
g++ graph_example.cpp ../graph.cpp -o graph_example -std=c++20
```

will do the same thing.

It is so annoying to enter such a long command every time! You can press `↑` to obtain your command history.

## VSCode configurations

Read [VSCodeConfig.md](#) if you need.

## Submission

---

Create a zip file `submission.zip` containing `graph.cpp` and `solve.cpp`, and then submit `submission.zip` to the OJ. You can `cd` to `attachments`, and then run the following command to create it:

```
zip -r submission.zip graph.cpp solve.cpp
```

Note that the files `graph.cpp` and `solve.cpp` should be placed directly in the zip file.